

NAREDDY JASHWANTH REDDY

Melbourne, VIC | +61 432 396 557 | jashwanthreddysungjin@gmail.com
[jashwanth-ai-portfolio.vercel.app](#) | [GitHub: NaReddyJashwanthReddy](#) | [Kaggle: jashwanthreddy0264](#)

SUMMARY

AI engineer with a year of production computer vision experience, now pursuing a Master of Artificial Intelligence at Monash. At Fotos, I connected image analysis with photo culling, editing and cloud inference. I am building on that experience through multi-agent systems, multimodal models and language-model adaptation.

EDUCATION

Monash University <i>Master of Artificial Intelligence</i>	2026 - Present Current WAM: 79.0 GPA: 3.5
BVRIT (JNTUH) <i>B.Tech in Electrical and Electronic Engineering</i> <i>Minor: Computer Science - Artificial Intelligence & Machine Learning</i>	2020 - 2024 CGPA: 8.05/10 CGPA: 8.90/10

EXPERIENCE

AI Engineer <i>Fotos - AI-powered photography workflows</i>	Feb 2025 - Jan 2026 Remote
- Deployed computer vision microservices on GCP using TensorFlow, OpenCV, RetinaFace and DeepFace/ArcFace to support face detection, facial analysis and photographic workflows. - Extended photo culling beyond face detection by combining identity clustering, quality/expression scoring, and Qwen2.5-VL served with vLLM, enabling context-aware image assessment and ranking. - Reduced manual editing time by 60% with neural preset style transfer; managed deployment, inference APIs, and monitoring.	
Artificial Intelligence Intern <i>Navodita InfoTech</i>	Jan 2024 - Feb 2024 Remote
Built a spaCy/NLTK chatbot with transformer-generated responses and an encoder-decoder French-English translator, connecting language-model experiments with practical conversation and translation tasks.	

SELECTED PROJECTS

AgentForge - In Development <i>Multi-Agent Systems</i> <i>Python, LangGraph, Groq</i> Details	
- Building a workspace that coordinates eight specialised roles across planning, research, product, marketing, development, testing, review and task management. - Defined task dependencies, shared project state and structured artifacts so specialists receive role-relevant context and hand work between stages. - Designed separate testing and requirements-review stages, revision loops and human approval gates; the portfolio presents the role architecture while implementation continues.	
Multimodal Image Analysis <i>Computer Vision</i> <i>Qwen2.5-VL, CLIP, YOLO, OpenCV</i> Details	
- Built an image-plus-text quality-assessment pipeline returning structured JSON with quality scores, identified visual drawbacks, and explanations, making model outputs easier to integrate with downstream applications. - Prototyped 4-bit Qwen2.5-VL-3B image-and-text inference in Colab; ran separate CLIP ranking and YOLO/OpenCV experiments for semantic matching and object-level analysis.	
LLM Fine-Tuning & Domain Adaptation <i>NLP</i> <i>Qwen3, Hugging Face, PyTorch</i> Details	
- Adapted Qwen3 in two stages: self-supervised continued pretraining on raw text from 20+ medical books and research papers, followed by supervised fine-tuning on labelled conversations. - Prepared role-structured, LLM-generated multi-turn examples and trained with a causal language-modeling objective for concise conversational replies; checked behaviour on unseen prompts. - Ran a separate PEFT query-rewriting study with 30 learned prompt tokens on 1,000 CANARD-derived examples (900 training / 100 validation), comparing adapted and base-model outputs.	

TECHNICAL SKILLS

Languages: Python, SQL.
ML & language: PyTorch, TensorFlow, Hugging Face, CNNs, transformers, fine-tuning, prompt tuning, byte-pair tokenization.
Vision & agents: OpenCV, RetinaFace, DeepFace/ArcFace, CLIP, YOLO, LangGraph, Groq.
Engineering: FastAPI, vLLM, GCP, Git/GitHub, REST APIs, deployment and monitoring.

ACHIEVEMENTS & CERTIFICATION

ATVC Innovation Runner-Up (2023) | My Anatomy AI-thon Finalist (2023)
Advanced Data Science & AI Certification - Intellipaat (IIT-Madras)